

Dynamic Programming

LQR471814

February 19, 2025

1 Basic ideas

Any dynamic programming problem has a certain objective, minimizing cost, maximizing profits, maximizing utility, etc... The function that describes this objective is known as the *objective function*.

Information about the current situation required to make a correct decision is known as *state*. (ex. To gauge how much to spend, you must know how much you have. Therefore wealth W would be one of the state variables)

Variables chosen by the agent are called *chosen variables*, these are decisions made based on *state variables* which influence the next state (though next state is often also influenced by stuff outside the *chosen variables*).

Dynamic programming describes the optimal chain of decisions by finding a rule that determines the most optimal decision given any possible of the state. (ex. If consumption only depends on wealth, you would want to find a function $c(W)$ which gives the optimal consumption for any degree of wealth) This is known as a *policy function*.

2 Formalization

Suppose x_t is the state at a time t , the initial state being x_0 .

At any point the set of possible actions is defined as $a_t \in \Gamma(x_t)$, where: 1) a_t are the particular values for a set of control variables (variables that influence the next state) 2) $\Gamma(x_t)$ is the set of actions that can be taken at state x_t .

State changes from x to a new state $T(x, a)$ when action a is taken.

The objective function that calculates the optimality of a given action a on a state x is $F(x, a)$.

Finally, we can define the *value function* (the function that describes the best possible value of the objective as a function of state x) like follows.

$$V(x) = \max_{a \in \Gamma(x)} \{F(x, a) + V(T(x, a))\}$$

Interpretation: The value function is the maximum value of the objective function (given the current a) added to the value function of the next state (given the current a) for all a in $\Gamma(x)$.

Usually we are given x_0 , T , F , and Γ as part of the problem description and are tasked with finding $V(x)$ (the most optimal value possible for a given starting state) and $a(x)$ (the optimal decision to make at a given state).

3 Examples

3.1 Optimal stopping problem

Suppose a person is evaluating potential employment opportunities for the next 10 years ($t = 1, 2, 3, \dots$).

At each value t , they may encounter a choice between a “good” job offering \$100 in salary or a “bad” job offering \$44, each with an equal probability of being offered.

- Pick the available job offered.
- Reject the offer and wait till the next year.

The question then becomes, what are the choices that should be made over the 10 year period to maximize money obtained?

This can be solved by reasoning backwards from $t = 10$:

1. At $t = 10$ the total earnings from accepting a “good” job is \$100; the value of accepting a “bad” job is \$44; the total earnings from rejecting

the available job is \$0. Therefore, if they are still unemployed in the last period, they should accept whatever job they are offered at that time for greater income.

2. At $t = 9$, the total earnings from accepting a “good” job is $2 \times 100 = 200$ because that job will last for two years. The total earnings from accepting a “bad” job is $2 \times 44 = 88$. The total expected earnings from rejecting a job offer are 0 now plus the value of the next job offer, which will either be 44 with $\frac{1}{2}$ probability or 100 with $\frac{1}{2}$ probability, for an average (“expected”) value of $\frac{\$100 + \$44}{2} = \$72$. Therefore, the job available at $t = 9$ should be accepted.
3. At $t = 8$, the total earnings from accepting a “good” job is $3 \times \$100 = \300 ; the total earnings from accepting a “bad” job is $3 \times \$44 = \132 . The total expected earnings from rejecting a job offer is \$0 now plus the total expected earnings from waiting for a job offer at $t = 9$. As previously concluded, any offer at $t = 9$ should be accepted and the expected value of doing so is $\frac{\$200 + \$88}{2} = \$144$. Therefore, at $t = 8$, total expected earnings are higher if the person waits for the next offer rather than accepting a “bad” job.

3.1.1 Math interpretation

$$\begin{aligned}
 x &= (o, j, t) \\
 x_o &\in \{44, 100\} \\
 x_j &\in \{\text{has job, no job}\} \\
 \{x_t &\in \mathbb{Z} | 1 \leq x \leq 10\}
 \end{aligned}$$

- x is the state.
- x_o is the amount of money you would earn if you had a job.
- x_j is your employment status.
- x_t is the current time.

$$a(x) \in \Gamma(x) = \begin{cases} \{\text{no change}\}, & x_j = \text{has job} \\ \{\text{no change, take job}\}, & x_j = \text{no job} \end{cases}$$

- $a(x)$ is the optimal action to take for a given state x .
- $\Gamma(x)$ is the possible actions that can be taken for a given state x .

$$T(x, a) = \begin{cases} (x_o, \text{has job}, x_t + 1), & a = \text{take job} \\ \begin{cases} (\text{rand}(100, 44), x_j, x_t + 1), & x_j = \text{no job} \\ (x_o, x_j, x_t + 1), & x_j = \text{has job} \end{cases}, & a = \text{no change} \end{cases}$$

- $T(x, a)$ defines the next state for a given current state x and chosen action a .

$$F(x, a) = \begin{cases} x_o, & a = \text{take job} \\ \begin{cases} x_o, & x_j = \text{has job} \\ 0, & x_j = \text{no job} \end{cases}, & a = \text{no change} \end{cases}$$

- $F(x, a)$ is the objective function that determines how optimal taking the given action a on the current state x is.

$$V(x) = \max_{a \in \Gamma(x)} \{F(x, a) + V(T(x, a))\}$$

We can see that we are given x , $\Gamma(x)$, $T(x, a)$, and $F(x, a)$ from the problem. Our job now, is to determine what $V(x)$ and $a(x)$ are, most commonly through processes of induction.

3.1.2 Solving for the equations

We'll use the technique of backwards induction to derive the equations $V(x)$ and $a(x)$.

This involves starting from a base case, then branching off from that base case and seeing what the commonalities are.

Base cases

$$\begin{aligned}
V((100, \text{no job}, 10)) &= \max\{100, 0\} \\
&= 100 \\
a((100, \text{no job}, 10)) &= \text{take job} \\
V((44, \text{no job}, 10)) &= \max\{44, 0\} \\
&= 44 \\
a((44, \text{no job}, 10)) &= \text{take job}
\end{aligned}$$

Since the offer has a probability of being either 100 or 44, we compute the “expected” optimal value at year 10 as the average between the optimal value at $x_o = 100$ or $x_o = 44$.

$$\begin{aligned}
V(\{100, 44\}, \text{no job}, 10) &= \frac{100 + 44}{2} \\
&= 72
\end{aligned}$$

If you have a job at year 10 then you just keep it (no other choices are available).

$$\begin{aligned}
V(\mathbb{R}, \text{has job}, 10) &= \max\{x_o\} \\
&= x_o \\
a(\mathbb{R}, \text{has job}, 10) &= \text{no change}
\end{aligned}$$

Starting from year 9

$$\begin{aligned}
V(\mathbb{R}, \text{has job}, 9) &= F(x, \text{no change}) + V(T(x, \text{no change})) \\
&= x_o + V((x_o, \text{has job}, 10)) \\
&= x_o + x_o \\
&= 2x_o
\end{aligned}$$

Like always, if you have a job already, you just keep it.

$$\begin{aligned}
V((100, \text{no job}, 9)) &= \max \begin{cases} F(x, \text{take job}) + V(T(x, \text{take job})) \\ F(x, \text{no change}) + V(T(x, \text{no change})) \end{cases} \\
&= \max\{72, 200\} \\
&= 200
\end{aligned}$$

$$\begin{aligned}
V((44, \text{no job}, 9)) &= \max \begin{cases} F(x, \text{take job}) + V(T(x, \text{take job})) \\ F(x, \text{no change}) + V(T(x, \text{no change})) \end{cases} \\
&= \max\{72, 88\} \\
&= 88
\end{aligned}$$

$$a(\{44, 100\}, \text{no job}, 9) = \text{take job}$$

And the expected optimal value at year 9.

$$V(\{44, 100\}, \text{no job}, 9) = 144$$

Starting from year 8

$$\begin{aligned}
V((\mathbb{R}, \text{has job}, 8)) &= F(x, \text{no change}) + V(T(x, \text{no change})) \\
&= 3x_o
\end{aligned}$$

$$\begin{aligned}
V((100, \text{no job}, 8)) &= \max \begin{cases} F(x, \text{take job}) + V(T(x, \text{take job})) \\ F(x, \text{no change}) + V(T(x, \text{no change})) \end{cases} \\
&= \max\{144, 300\} \\
&= 300
\end{aligned}$$

$$\begin{aligned}
V((44, \text{no job}, 8)) &= \max \begin{cases} F(x, \text{take job}) + V(T(x, \text{take job})) \\ F(x, \text{no change}) + V(T(x, \text{no change})) \end{cases} \\
&= \max\{144, 132\} \\
&= 144
\end{aligned}$$

$$a((100, \text{no job}, 8)) = \text{take job}$$

$$a((44, \text{no job}, 8)) = \text{no change}$$

$$\begin{aligned}
V(\{44, 100\}, \text{no job}, 8) &= \frac{300 + 144}{2} \\
&= 222
\end{aligned}$$

So by thinking very hard now, we can see that the optimal strategy is:

$$a(x) = \begin{cases} \text{no change,} & x_j = \text{has job} \\ \left\{ \begin{array}{l} \text{take job,} \\ \text{no change,} \end{array} \right. & \begin{array}{l} x_t \geq 9 \\ x_o = 100 \\ x_o = 44 \end{array}, \quad x_t < 9, \quad x_j = \text{no job} \end{cases}$$